

Problem Statement: Abalone Age Prediction

Description:- Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Attribute Information:

Given is the attribute name, attribute type, measurement unit, and a brief description. The number of rings is the value to predict: either as a continuous value or as a classification problem.

Name/Data Type / Measurement Unit / Description

- 1- Sex / nominal /--/M, F, and I (infant)
- 2- Length/continuous/mm/Longest shell measurement
- 3- Diameter / continuous/mm/perpendicular to length
- 4- Height / continuous/mm/with meat in shell
- 5- Whole weight/continuous/grams/whole abalone
- 6- Shucked weight/continuous/grams / weight of meat 7- Viscera weight/continuous/grams/gut weight (after bleeding)
- 8- Shell weight/continuous/grams/ after being dried
- 9- Rings/integer /--/ +1.5 gives the age in years

Building a Regression Model

1. Download the dataset: Dataset
2. Load the dataset into the tool. 3. Perform Below Visualizations.

Univariate Analysis

Bi-Variate Analysis

Multi-Variate Analysis

4. Perform descriptive statistics on the dataset.
5. Check for Missing values and deal with them.
6. Find the outliers and replace them outliers
7. Check for Categorical columns and perform encoding. 8. Split the data into dependent and independent variables. 9. Scale the independent variables
10. Split the data into training and testing
11. Build the Model 12. Train the Model
13. Test the Model
14. Measure the performance using Metrics.

Data Description¶¶

Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict the age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

From the original data examples with missing values were removed (the majority having the predicted value missing), and the ranges of the continuous values have been scaled for use with an ANN (by dividing by 200).

Attribute Information:

Given is the attribute name, attribute type, the measurement unit and a brief description. The number of rings is the value to predict: either as a continuous value or as a classification problem.

Name / Data Type / Measurement Unit / Description
Sex / nominal / -- / M, F, and I (infant)

Length / continuous / mm / Longest shell measurement

Diameter / continuous / mm / perpendicular to length

Height / continuous / mm / with meat in shell

Whole weight / continuous / grams / whole abalone

Shucked weight / continuous / grams / weight of meat

Viscera weight / continuous / grams / gut weight (after bleeding)

Shell weight / continuous / grams / after being dried

Rings / integer / -- / +1.5 gives the age in years

```
import libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
import seaborn as sns
```

```
df = pd.read_csv('../input/abalone.csv')
```

```
df.head()
```

Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
-----	--------	----------	--------	--------------	----------------	----------------	--------------	-------

0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
---	---	-------	-------	-------	--------	--------	--------	-------	----

1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
---	---	-------	-------	-------	--------	--------	--------	-------	---

2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
---	---	-------	-------	-------	--------	--------	--------	-------	---

3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
---	---	-------	-------	-------	--------	--------	--------	-------	----

4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7
---	---	-------	-------	-------	--------	--------	--------	-------	---

```
df.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
--	--------	----------	--------	--------------	----------------	----------------	--------------	-------

count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
-------	-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
------	----------	----------	----------	----------	----------	----------	----------	----------

std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
-----	----------	----------	----------	----------	----------	----------	----------	----------

min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
-----	----------	----------	----------	----------	----------	----------	----------	----------

```

25%  0.450000    0.350000    0.115000    0.441500    0.186000    0.093500
0.130000    8.000000
50%  0.545000    0.425000    0.140000    0.799500    0.336000    0.171000
0.234000    9.000000
75%  0.615000    0.480000    0.165000    1.153000    0.502000    0.253000
0.329000    11.000000
max   0.815000    0.650000    1.130000    2.825500    1.488000    0.760000
1.005000    29.000000

```

```
df['age'] = df['Rings']+1.5
```

```
df = df.drop('Rings', axis = 1)
```

```
EDA
```

```
sns.heatmap(df.isnull())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcc468da358>
```

```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7fcc3caa8160>
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

```
Sex          4177 non-null object
```

```
Length       4177 non-null float64
```

```
Diameter     4177 non-null float64
```

```
Height       4177 non-null float64
```

```
Whole weight 4177 non-null float64
```

```
Shucked weight 4177 non-null float64
```

```
Viscera weight 4177 non-null float64
```

```
Shell weight  4177 non-null float64
```

```
age           4177 non-null float64
```

```
dtypes: float64(8), object(1)
```

```
memory usage: 293.8+ KB
```

```
numerical_features = df.select_dtypes(include = [np.number]).columns
```

```
categorical_features = df.select_dtypes(include = [np.object]).columns
```

```
numerical_features
```

```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'age'],
      dtype='object')
```

```
categorical_features
```

```
Index(['Sex'], dtype='object')
```

```
plt.figure(figsize = (20,7))
```

```
sns.heatmap(df[numerical_features].corr(),annot = True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcc29714dd8>
```

Whole Weight is almost linearly varying with all other features except age
Height has least linearity with remaining features
Age is most linearly proportional with Shell Weight followed by Diameter and length
Age is least correlated with Shucked Weight
Key insight:

All numerical features but 'sex'

- Though features are not normally distributed, are close to normality
- None of the features have minimum = 0 except Height (requires re-check)
- Each feature has difference scale range

```
sns.countplot(x = 'Sex', data = df, palette = 'Set3')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcc26ba6748>
```

```
plt.figure(figsize = (20,7))
```

```
sns.swarmplot(x = 'Sex', y = 'age', data = df, hue = 'Sex')
```

```
sns.violinplot(x = 'Sex', y = 'age', data = df)
```

```
/opt/conda/lib/python3.6/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.
```

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcc26b67b38>
```

Male : age majority lies in between 7.5 years to 19 years

Female: age majority lies in between 8 years to 19 years

Immature: age majority lies in between 6 years to < 10 years

Data Preprocessing

outlier handling

```
df = pd.get_dummies(df)
```

```
dummy_df = df
```

```
var = 'Viscera weight'
```

```
plt.scatter(x = df[var], y = df['age'])
```

```
plt.grid(True)
```

```
df.drop(df[(df['Viscera weight'] > 0.5) &  
(df['age'] < 20)].index, inplace = True)
```

```
df.drop(df[(df['Viscera weight'] < 0.5) & (  
df['age'] > 25)].index, inplace = True)
```

```
var = 'Shell weight'
```

```
plt.scatter(x = df[var], y = df['age'])
```

```
plt.grid(True)
```

```
df.drop(df[(df['Shell weight'] > 0.6) &  
(df['age'] < 25)].index, inplace = True)
```

```
df.drop(df[(df['Shell weight']<0.8) & (  
df['age'] > 25)].index, inplace = True)  
var = 'Shucked weight'  
plt.scatter(x = df[var], y = df['age'])  
plt.grid(True)
```

```
df.drop(df[(df['Shucked weight'] >= 1) &  
(df['age'] < 20)].index, inplace = True)  
df.drop(df[(df['Viscera weight']<1) & (  
df['age'] > 20)].index, inplace = True)  
var = 'Whole weight'  
plt.scatter(x = df[var], y = df['age'])  
plt.grid(True)
```

```
df.drop(df[(df['Whole weight'] >= 2.5) &  
(df['age'] < 25)].index, inplace = True)  
df.drop(df[(df['Whole weight']<2.5) & (  
df['age'] > 25)].index, inplace = True)  
var = 'Diameter'  
plt.scatter(x = df[var], y = df['age'])  
plt.grid(True)
```

```
df.drop(df[(df['Diameter'] <0.1) &  
(df['age'] < 5)].index, inplace = True)  
df.drop(df[(df['Diameter']<0.6) & (  
df['age'] > 25)].index, inplace = True)  
df.drop(df[(df['Diameter']>=0.6) & (  
df['age'] < 25)].index, inplace = True)  
var = 'Height'  
plt.scatter(x = df[var], y = df['age'])  
plt.grid(True)
```

```
df.drop(df[(df['Height'] > 0.4) &  
(df['age'] < 15)].index, inplace = True)  
df.drop(df[(df['Height']<0.4) & (  
df['age'] > 25)].index, inplace = True)  
var = 'Length'  
plt.scatter(x = df[var], y = df['age'])  
plt.grid(True)
```

```
df.drop(df[(df['Length'] <0.1) &  
(df['age'] < 5)].index, inplace = True)  
df.drop(df[(df['Length']<0.8) & (  
df['age'] > 25)].index, inplace = True)
```

```

df.drop(df[(df['Length']>=0.8) & (
df['age'] < 25)].index, inplace = True)
Feature Selection and Standardization
X = df.drop('age', axis = 1)
y = df['age']
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
standardScale = StandardScaler()
standardScale.fit_transform(X)

selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)

X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)
/opt/conda/lib/python3.6/site-packages/sklearn/preprocessing/data.py:645:
DataConversionWarning: Data with input dtype uint8, float64 were all converted to float64 by
StandardScaler.
    return self.partial_fit(X, y)
/opt/conda/lib/python3.6/site-packages/sklearn/base.py:464: DataConversionWarning: Data with
input dtype uint8, float64 were all converted to float64 by StandardScaler.
    return self.fit(X, **fit_params).transform(X)
Model Selection
1)Linear regression
from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train, y_train)
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
y_train_pred = lm.predict(X_train)
y_test_pred = lm.predict(X_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error
s = mean_squared_error(y_train, y_train_pred)
print('Mean Squared error of training set :%2f%s)

p = mean_squared_error(y_test, y_test_pred)
print('Mean Squared error of testing set :%2f%p)
Mean Squared error of training set :3.551893
Mean Squared error of testing set :3.577687
from sklearn.metrics import r2_score
s = r2_score(y_train, y_train_pred)
print('R2 Score of training set:%.2f%s)

p = r2_score(y_test, y_test_pred)

```

```
print('R2 Score of testing set:%.2f%p)
```

```
R2 Score of training set:0.54
```

```
R2 Score of testing set:0.53
```

```
2)Ridge
```

```
from sklearn.linear_model import Ridge
```

```
ridge_mod = Ridge(alpha=0.01, normalize=True)
```

```
ridge_mod.fit(X_train, y_train)
```

```
ridge_mod.fit(X_test, y_test)
```

```
ridge_model_pred = ridge_mod.predict(X_test)
```

```
ridge_mod.score(X_train, y_train)
```

```
0.5307346478347332
```

```
ridge_mod.score(X_test, y_test)
```

```
0.5272608729607438
```

```
plt.scatter(y_test, ridge_model_pred)
```

```
plt.xlabel('True Values')
```

```
plt.ylabel('Predictions')
```

```
Text(0, 0.5, 'Predictions')
```

```
3)Support vector Regression
```

```
from sklearn.svm import SVR
```

```
# LINEAR KERNEL
```

```
svr = SVR(kernel = 'linear')
```

```
svr.fit(X_train, y_train)
```

```
svr.fit(X_test, y_test)
```

```
SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
```

```
    gamma='auto_deprecated', kernel='linear', max_iter=-1, shrinking=True,
```

```
    tol=0.001, verbose=False)
```

```
y_train_pred = svr.predict(X_train)
```

```
y_test_pred = svr.predict(X_test)
```

```
svr.score(X_train, y_train)
```

```
0.4461014542389635
```

```
svr.score(X_test, y_test)
```

```
0.43681391121982105
```

```
4) RandomForestRegression
```

```
from sklearn.ensemble import RandomForestRegressor
```

```
regr = RandomForestRegressor(max_depth=2, random_state=0,
```

```
                             n_estimators=100)
```

```
regr.fit(X_train, y_train)
```

```
regr.fit(X_test, y_test)
```

```
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=2,
```

```
                       max_features='auto', max_leaf_nodes=None,
```

```
                       min_impurity_decrease=0.0, min_impurity_split=None,
```



```
min_samples_leaf=1, min_samples_split=2,  
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None,  
oob_score=False, random_state=0, verbose=0, warm_start=False)  
y_train_pred = regr.predict(X_train)  
y_test_pred = regr.predict(X_test)
```

```
regr.score(X_train, y_train)  
0.4287379777803546
```

```
regr.score(X_test, y_test)  
0.43753106247261264
```

5)Gradient Boosting Regressor

```
from sklearn.ensemble import GradientBoostingRegressor
```

```
gbr = GradientBoostingRegressor()
```

```
gbr.fit(X_train, y_train)
```

```
gbr.fit(X_test, y_test)
```

```
GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,  
learning_rate=0.1, loss='ls', max_depth=3, max_features=None,  
max_leaf_nodes=None, min_impurity_decrease=0.0,  
min_impurity_split=None, min_samples_leaf=1,  
min_samples_split=2, min_weight_fraction_leaf=0.0,  
n_estimators=100, n_iter_no_change=None, presort='auto',  
random_state=None, subsample=1.0, tol=0.0001,  
validation_fraction=0.1, verbose=0, warm_start=False)
```

```
y_train_pred = regr.predict(X_train)
```

```
y_test_pred = regr.predict(X_test)
```

```
regr.score(X_train, y_train)
```

```
0.4287379777803546
```

```
regr.score(X_test, y_test)
```

```
0.43753106247261264
```

6)KNeighborsRegressor

```
from sklearn.neighbors import KNeighborsRegressor
```

```
knn = KNeighborsRegressor(n_neighbors =4 )
```

```
knn.fit(X_train, y_train)
```

```
knn.fit(X_test, y_test)
```

```
KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',  
metric_params=None, n_jobs=None, n_neighbors=4, p=2,  
weights='uniform')
```

```
y_train_pred = knn.predict(X_train)
```

```
y_test_pred = knn.predict(X_test)
```

```
knn.score(X_train, y_train)
```

```
0.4677575709446231
```

```
knn.score(X_test, y_test)
```

```
0.6856343678141352
```

you have seen the performance of each one of above model.

so according to you which model should we start or choose?

"Suppose there exist two explanations for an occurrence. In this case the simpler one is usually better. Another way of saying it is that the more assumptions you have to make, the more unlikely an explanation." Hence, starting with the simplest model Ridge, for various reasons:

- Feature Dimension is less
- No missing values
- Few categorical features

Hyperparameter Tunning Using GridSearchCV

Hyperparameter Tuning using GridSearchCV

```
from sklearn.model_selection import GridSearchCV
param = {'alpha':[0.01, 0.1, 1,10,100],
        'solver' : ['auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', 'saga']}
glrm0 = GridSearchCV(estimator = Ridge(random_state=10,),
param_grid = param,scoring= 'r2' ,cv = 5, n_jobs = -1)
glrm0.fit(X_train, y_train)
glrm0.best_params_, glrm0.best_score_
({'alpha': 0.1, 'solver': 'sag'}, 0.5308712206007367)
ridge_mod = Ridge(alpha=0.001,solver = 'sag', random_state = 10, normalize=True)
ridge_mod.fit(X_train, y_train)
ridge_mod.fit(X_test, y_test)
ridge_model_pred = ridge_mod.predict(X_test)
ridge_mod.score(X_train, y_train)
0.5331463016536355
ridge_mod.score(X_test, y_test)
0.534651599310652
```